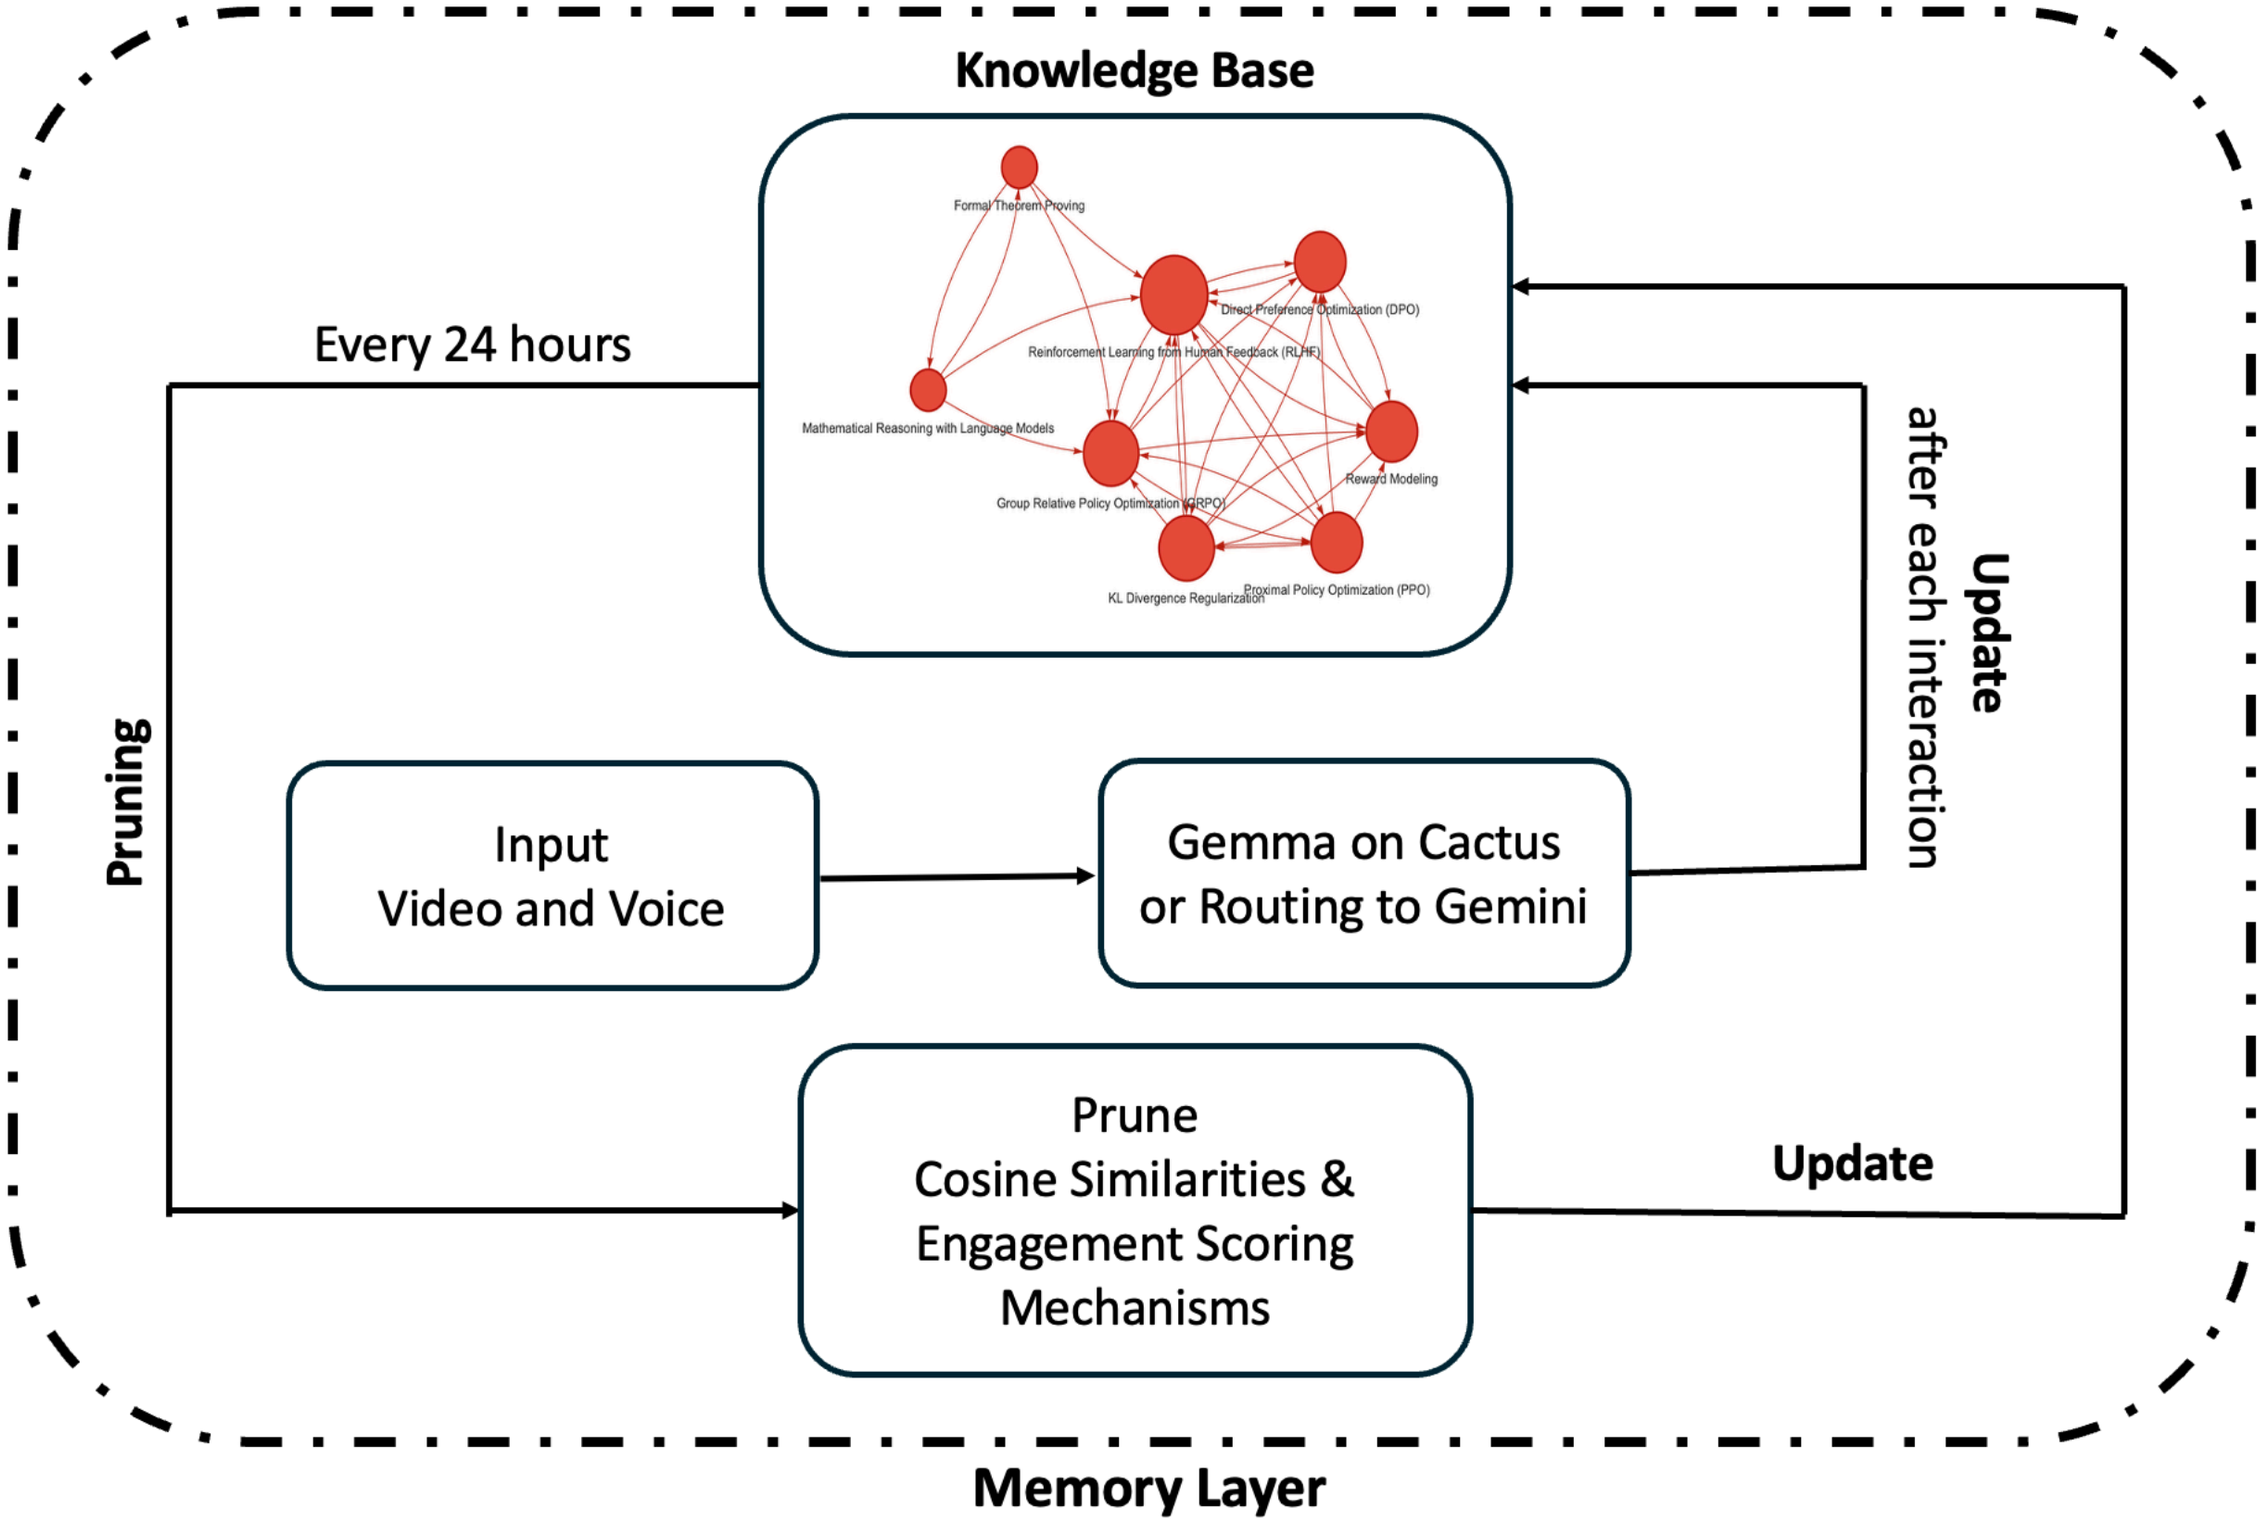
Knowledge Base
Formal Theorem Proving
Direct Preference Optimization (DPO)
Reinforcement Learning from Human Feedback (RLHF)
Mathematical Reasoning with Language Models
Reward Modeling
Group Relative Policy Optimization (GRPO)
KL Divergence Regularization
Proximal Policy Optimization (PPO)
Every 24 hours
Pruning
Update
after each interaction
Input
Video and Voice
Gemma on Cactus
or Routing to Gemini
Prune
Cosine Similarities &
Engagement Scoring
Mechanisms
Update
Memory Layer